

Retrieving Wikipedia articles

In this module, we focused on using nearest neighbors and clustering to retrieve documents that interest users, by analyzing their text. We explored two document representations: word counts and TF-IDF. We also built an Jupyter notebook for retrieving articles from Wikipedia about famous people.

In this assignment, we are going to dig deeper into this application, explore the retrieval results for various famous people, and familiarize ourselves with the code needed to build a retrieval system.

Follow the rest of the instructions on this page to complete your program. When you are done, *you will answer a series of questions*. The instructions will indicate what data to collect for answering the quiz.

Learning outcomes

- Execute document retrieval code with the iPython notebook
- Load and transform real, text data
- Compare results with word counts and TF-IDF
- Set the distance function in the retrieval
- Build a document retrieval model using nearest neighbor search

Download the data and starter code

Before getting started, you will need to download the dataset and the starter iPython notebook that we used in the module.

- Download the wikipedia dataset with articles on famous people here in SFrame format: *people_wiki.gl.zip*
- Download the document retrieval notebook from the module here: *document-retrieval.ipynb*
- Save both of these files in the same directory (where you are calling Jupyter notebook from) and unzip the data file.

Now you are ready to get started!

What you will do

We are going to do three tasks in this assignment. There are several results you need to gather along the way to enter into the quiz after this reading.

1. **Compare top words according to word counts to TF-IDF:** In the notebook we covered in the module, we explored two document representations: word counts and TF-IDF. Now, take a particular famous person, 'Elton John'. **Questions 1: What are the 3 words in his articles with highest word counts? What are the 3 words in his articles with highest TF-IDF?** These results illustrate why TF-IDF is useful for finding important words.
2. **Measuring distance:** Elton John is a famous singer; let's compute the distance between his article and those of two other famous singers. In this assignment, you will use the *cosine distance*, which is one measure of similarity between vectors, similar to the one discussed in the lectures. You can compute this distance using the `graphlab.distances.cosine` function. **Questions 2: What's the cosine distance between the articles on 'Elton John' and 'Victoria Beckham'? What's the cosine distance between the articles on 'Elton John' and 'Paul McCartney'? Which one of the two is closest to Elton John? Does this result make sense to you?**
3. **Building nearest neighbors models with different input features and setting the distance metric:** In the sample notebook, we built a nearest neighbors model for retrieving articles using TF-IDF as features and using the default setting in the construction of the nearest neighbors model. Now, you will build two nearest neighbors models:
 - Using word counts as features
 - Using TF-IDF as features

In both of these models, we are going to set the distance function to cosine similarity.

Now we are ready to use our model to retrieve documents. Use these two models to collect the following results, **Questions 3:**

- **What's the most similar article, other than itself, to the one on 'Elton John' using word count features?**
- **What's the most similar article, other than itself, to the one on 'Elton John' using TF-IDF features?**
- **What's the most similar article, other than itself, to the one on 'Victoria Beckham' using word count features?**
- **What's the most similar article, other than itself, to the one on 'Victoria Beckham' using TF-IDF features?**